
The Simpsons Characters

—— Sijie Huang, Hanying Shen, Yijia Sun ——

1. Ask a Question

- Simpson character image dataset
- To distinguish the characters of Homer Simpson and Lisa Simpson
- Use CNN and Transfer Learning to do the binary image classification



2. Get the data.

- Kaggle: <https://www.kaggle.com/alexattia/the-simpsons-characters-dataset>
- 42 classes with 20933 images

- Homer Simpson: 2247 images
- Lisa Simpson: 1354 images

- Data cleaning: add labels to images before training the model
- Related dataset: test set (to test the model at the end)

3. Explore the Data

- Problem with data?
 - The original dataset has 42 classes, and many of them have less than 100 images
 - Too many classes will lead to less accuracy
- How to improve this dataset in the future?
 - Capture more images
 - Reduce to less classes

36	Fat Tony	27	23	4	0
37	Gil	27	23	4	0
38	Miss Hoov	17	14	3	0
39	Disco Stu	8	7	1	0
40	Troy Mccl	8	7	1	0
41	Lionel Hut	3	3	0	0
42	Jimbo Jon	0	0	0	0
43	Bumblebe	0	0	0	0
44	Hans Mole	0	0	0	0
45	Helen Lov	0	0	0	0
46	Jasper Bea	0	0	0	0

4. Model the Data

- Image dataset → CNN and Transfer Learning, SVM
- Accuracy: Training and Validation Loss, Confusion Matrix, ROC

5. Select a model and train it.

- Random Forest Classifier
- SVM (Support Vector Machine)
- CNN
- Transfer Learning


```
precision = accuracy_score(test_predictions, y_test) * 100  
print("Accuracy with RandomForest: {0:.6f}".format(precision))
```

Accuracy with RandomForest: 77.222222

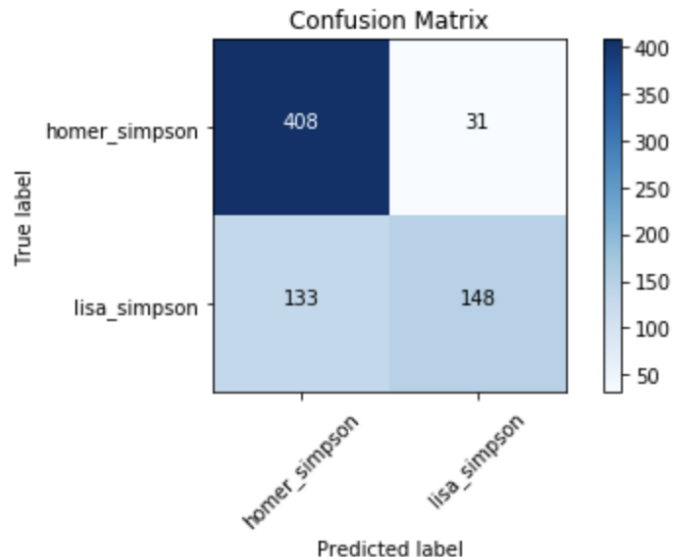
Random Forest Classifier

Confusion matrix, without normalization

```
[[408  31]  
 [133 148]]
```




AUC test score: 0.73




```
precision = accuracy_score(test_predictions1, y_test) * 100  
print("Accuracy with SVM: {0:.6f}".format(precision))
```

Accuracy with SVM: 77.222222



Confusion matrix, without normalization

```
[[380  59]  
 [105 176]]
```



AUC test score: 0.75

CNN

```
model = Sequential([
    Conv2D(16, 3, padding='same', activation='relu', input_shape=(224, 224 ,3)
    MaxPooling2D(),
    Conv2D(32, 3, padding='same', activation='relu'),
    MaxPooling2D(),
    Conv2D(64, 3, padding='same', activation='relu'),
    MaxPooling2D(),
    Flatten(),
    Dense(512, activation='relu'),
    Dense(2, activation='sigmoid')
])
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

Epoch 1/5

4/4 [=====] - 37s 9s/step - loss: 0.7097 - accuracy:
0.6150 - val_loss: 0.6741 - val_accuracy: 0.6233

Epoch 2/5

4/4 [=====] - 33s 8s/step - loss: 0.6748 - accuracy:
0.6091 - val_loss: 0.6702 - val_accuracy: 0.6409

Epoch 3/5



4/4 [=====] - 36s 9s/step - loss: 0.6639 - accuracy:
0.6360 - val_loss: 0.6525 - val_accuracy: 0.6178



Epoch 4/5



4/4 [=====] - 34s 9s/step - loss: 0.6570 - accuracy:
0.6149 - val_loss: 0.6518 - val_accuracy: 0.6139

Epoch 5/5

4/4 [=====] - 35s 9s/step - loss: 0.6425 - accuracy:
0.6271 - val_loss: 0.6494 - val_accuracy: 0.6100







6. Fine-tune the model.

Import VGG16 model

```
#Build Fine-tuned VGG16 model
```

Freeze all layers besides
`predictions`

```
vgg16_model = keras.applications.vgg16.VGG16()
```

```
model = Sequential()  
for layer in vgg16_model.layers:  
    if layer.name != 'predictions':  
        model.add(layer)
```

Add our layer to the model

```
for layer in model.layers:  
    layer.trainable = False
```

```
model.add(Dense(2, activation='softmax'))
```

Epoch 1/5

- 547s - loss: 0.6805 - accuracy: 0.5849 - val_loss: 0.6583 - val_accuracy: 0.6150

Epoch 2/5

- 534s - loss: 0.6640 - accuracy: 0.6206 - val_loss: 0.6908 - val_accuracy: 0.6101

Epoch 3/5

- 645s - loss: 0.6486 - accuracy: 0.6284 - val_loss: 0.6487 - val_accuracy: 0.6300


Epoch 4/5

- 505s - loss: 0.6334 - accuracy: 0.6471 - val_loss: 0.6131 - val_accuracy: 0.6352


Epoch 5/5

- 554s - loss: 0.6206 - accuracy: 0.6277 - val_loss: 0.5780 - val_accuracy: 0.6750

Training and Validation Accuracy

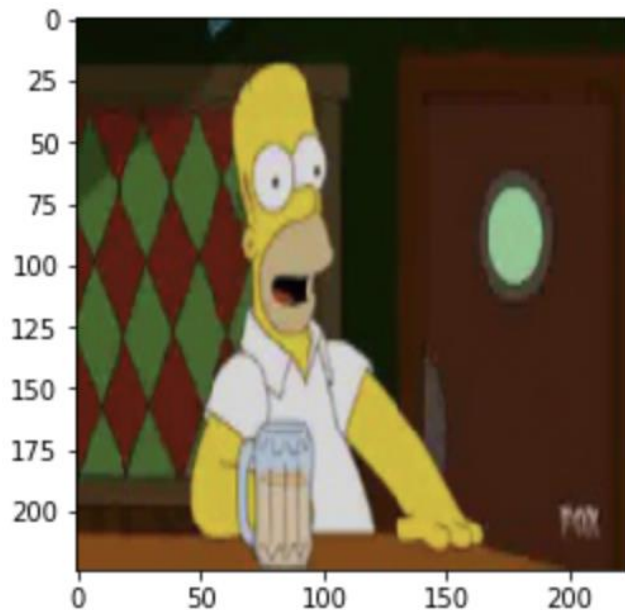


Training and Validation Loss



Solution

- Upload an image
- Resize it
- Use the model to predict who is it in the image



Most likely class: homer_simpson -- Probability: 0.9114153
Most likely class: lisa_simpson -- Probability: 0.09846259